

NAS Exploration

November 4, 2017

1 Gramener Data Science Entry Level Position Use Case 1

1.1 Introduction

For use case 1, we're examining the National Achievement Survey for class VIII students from 2014. In an effort to answer the following questions - 1. What influences student performance the most? 2. How do girls and boys perform across states? 3. Do students from South India really excel at Math and Science

```
In [16]: #imports
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

In [17]: datapath = './'
filename = 'nas-pupil-marks.csv'
Data = pd.read_csv(datapath+filename)
Data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 185348 entries, 0 to 185347
Data columns (total 64 columns):
STUID                185348 non-null int64
State                185348 non-null object
District             185348 non-null int64
Gender               185348 non-null int64
Age                  185348 non-null int64
Category             185348 non-null int64
Same language        185348 non-null int64
Siblings              185348 non-null int64
Handicap              185348 non-null int64
Father edu            185348 non-null int64
Mother edu           185348 non-null int64
Father occupation     185348 non-null int64
Mother occupation     185348 non-null int64
Below poverty         185348 non-null int64
Use calculator        185348 non-null int64
```

Use computer	166186	non-null	object
Use Internet	185348	non-null	int64
Use dictionary	185348	non-null	int64
Read other books	185348	non-null	int64
# Books	185348	non-null	int64
Distance	185348	non-null	int64
Computer use	185348	non-null	int64
Library use	185348	non-null	int64
Like school	185348	non-null	int64
Subjects	185348	non-null	object
Give Lang HW	185348	non-null	int64
Give Math HW	185348	non-null	int64
Give Scie HW	185348	non-null	int64
Give SoSc HW	185348	non-null	int64
Correct Lang HW	185348	non-null	int64
Correct Math HW	185348	non-null	int64
Correct Scie HW	185348	non-null	int64
Correct SocS HW	185348	non-null	int64
Help in Study	185348	non-null	int64
Private tuition	185348	non-null	int64
English is difficult	185348	non-null	int64
Read English	185348	non-null	int64
Dictionary to learn	185348	non-null	int64
Answer English WB	185348	non-null	int64
Answer English aloud	185348	non-null	int64
Maths is difficult	185348	non-null	int64
Solve Maths	185348	non-null	int64
Solve Maths in groups	185348	non-null	int64
Draw geometry	185348	non-null	int64
Explain answers	185348	non-null	int64
SocSci is difficult	185348	non-null	int64
Historical excursions	185348	non-null	int64
Participate in SocSci	185348	non-null	int64
Small groups in SocSci	185348	non-null	int64
Express SocSci views	185348	non-null	int64
Science is difficult	185348	non-null	int64
Observe experiments	185348	non-null	int64
Conduct experiments	185348	non-null	int64
Solve science problems	185348	non-null	int64
Express science views	185348	non-null	int64
Watch TV	185348	non-null	int64
Read magazine	185348	non-null	int64
Read a book	185348	non-null	int64
Play games	185348	non-null	int64
Help in household	185348	non-null	int64
Maths %	92681	non-null	float64
Reading %	93271	non-null	float64
Science %	90992	non-null	float64

```
Social %                89571 non-null float64
dtypes: float64(4), int64(57), object(3)
memory usage: 90.5+ MB
```

1.2 Data Cleaning

From our info function, we can see there are a few null entries in our dataset, thus it becomes important to try and handle those first. The rows which have the most nulls are our metrics of performance, the maths, science, social science and reading scores.

```
In [18]: NullsDropped = Data.dropna(axis=0,how='all',inplace=False)
        NullsDropped.shape
```

```
Out[18]: (185348, 64)
```

Immediately, there are no rows such that all values are null. The fact that our nulls are in our target variables, presents a problem as we can't easily replace them with a statistical measure without tainting the spread of features.

Dropping all the rows with nulls would also not be ideal, as we'd reduce our current number of data points by ~96% as is evidenced below.

```
In [19]: NullsDropped = Data.dropna(axis=0,how='any',inplace=False)
        NullsDropped.shape
```

```
Out[19]: (8044, 64)
```

In order to preserve the majority of our dataset, we split our target variables - we separate our current dataset to 5 different ones, one each for our four performance metrics and one for Maths and Science.

Moreover, our dataset also has zeroes in our dataset for columns that don't have valid 0 levels. This also needs to be fixed.

```
In [20]: Maths = Data.drop(['Reading %','Science %','Social %'],axis=1)
```

Similarly, we create a dataset for each metric of performance

```
In [21]: Reading = Data.drop(['Maths %','Science %','Social %'],axis=1)
        Science = Data.drop(['Maths %','Reading %','Social %'],axis=1)
        Social = Data.drop(['Maths %','Reading %','Science %'],axis=1)
        Maths_and_Science = Data.drop(['Reading %','Social %'],axis=1)
```

```
In [22]: filename = 'nas-labels.csv'
        labels = pd.read_csv(datapath+filename)
        labels.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 241 entries, 0 to 240
Data columns (total 4 columns):
Column      241 non-null object
```

```
Name      241 non-null object
Level     241 non-null object
Rename    241 non-null object
dtypes: object(4)
memory usage: 7.6+ KB
```

```
In [23]: # We use the labels.csv file to figure out which columns in our dataset are allowed to
```

```
zero_allowed = pd.DataFrame(labels.loc[labels['Level'].astype(str) == '0'])
```

```
#Slightly janky, but I didn't want to drop the states and can't convert characters to
```

```
zero_allowed = zero_allowed['Column'].tolist()
```

```
#Gives me the list of the columns which shouldn't have zeroes.
```

```
def remove_zeroes(df):
```

```
'''Helper function that first removes nulls and then removes 0's from rows that s
```

```
df.dropna(axis=0,how='any',inplace=True)
```

```
zero_not_allowed = [item for item in df.columns.values.tolist() if item not in zero
```

```
df = df.replace(0,np.nan)
```

```
#Since the dropna function has a subset parameter, casting 0's to np.nan
```

```
df.dropna(axis =0, how='any',inplace=True,subset = zero_not_allowed)
```

```
df.fillna(value=0,inplace=True)
```

```
#Convert object types to categories and categories to cat codes.
```

```
objects = df.select_dtypes(include=['O']).columns.tolist()
```

```
for col in objects:
```

```
    df[col] = df[col].astype('category')
```

```
    return df
```

```
In [24]: Maths=remove_zeroes(Maths)
```

```
Science=remove_zeroes(Science)
```

```
Social=remove_zeroes(Social)
```

```
Reading=remove_zeroes(Reading)
```

```
Maths_and_Science=remove_zeroes(Maths_and_Science)
```

```
NullsDropped=remove_zeroes(NullsDropped)
```

```
/home/karmanya/anaconda3/envs/MLIntel/lib/python3.5/site-packages/ipykernel/__main__.py:13: Set
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
```

Finally we should take the average of the performance metrics for the frames with multiple metrics

```

In [25]: ## Maths and Science
Maths_and_Science['Avg %'] = (Maths_and_Science['Maths %']+Maths_and_Science['Science %'])
Maths_and_Science.drop(['Maths %','Science %'],axis=1,inplace=True)

NullsDropped['Avg %'] = (NullsDropped['Maths %']+NullsDropped['Science %']+NullsDropped['Social %']+NullsDropped['Reading %'])
NullsDropped.drop(['Maths %','Science %','Social %','Reading %'],axis=1,inplace=True)

In [26]: Science.isnull().values.any()

Out[26]: False

In [27]: Maths.describe()

Out[27]:
```

	STUID	District	Gender	Age	Category \
count	5.054200e+04	50542.000000	50542.000000	50542.000000	50542.000000
mean	2.714830e+10	5.114222	1.520894	3.644870	2.654901
std	9.808741e+09	4.030791	0.499568	0.946113	1.211942
min	1.101100e+10	1.000000	1.000000	1.000000	0.000000
25%	1.808118e+10	2.000000	1.000000	3.000000	2.000000
50%	2.607110e+10	4.000000	2.000000	4.000000	3.000000
75%	3.601103e+10	7.000000	2.000000	4.000000	4.000000
max	4.502103e+10	28.000000	2.000000	6.000000	4.000000

	Same language	Siblings	Handicap	Father edu	Mother edu \
count	50542.000000	50542.000000	50542.000000	50542.000000	50542.000000
mean	1.305647	3.465039	1.895374	2.213090	1.885363
std	0.531592	1.209956	0.380640	1.222056	1.176629
min	0.000000	1.000000	0.000000	0.000000	0.000000
25%	1.000000	2.000000	2.000000	1.000000	1.000000
50%	1.000000	3.000000	2.000000	2.000000	2.000000
75%	2.000000	5.000000	2.000000	3.000000	3.000000
max	2.000000	5.000000	2.000000	5.000000	5.000000

	...	Observe experiments	Conduct experiments \
count	...	50542.000000	50542.000000
mean	...	2.646769	2.082941
std	...	0.692864	0.930929
min	...	1.000000	1.000000
25%	...	3.000000	1.000000
50%	...	3.000000	2.000000
75%	...	3.000000	3.000000
max	...	3.000000	3.000000

	Solve science problems	Express science views	Watch TV \
count	50542.000000	50542.000000	50542.000000
mean	2.426833	2.349828	3.556587
std	0.813011	0.822457	0.798879
min	1.000000	1.000000	1.000000
25%	2.000000	2.000000	3.000000

50%	3.000000	3.000000	4.000000
75%	3.000000	3.000000	4.000000
max	3.000000	3.000000	4.000000

	Read magazine	Read a book	Play games	Help in household \
count	50542.000000	50542.000000	50542.000000	50542.000000
mean	3.168850	3.096019	3.445333	3.511456
std	1.017827	0.922589	0.882576	1.012895
min	1.000000	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000	4.000000
50%	3.000000	3.000000	4.000000	4.000000
75%	4.000000	4.000000	4.000000	4.000000
max	4.000000	4.000000	4.000000	4.000000

	Maths %
count	50542.000000
mean	32.903478
std	15.862777
min	1.670000
25%	22.030000
50%	28.810000
75%	38.980000
max	100.000000

[8 rows x 58 columns]

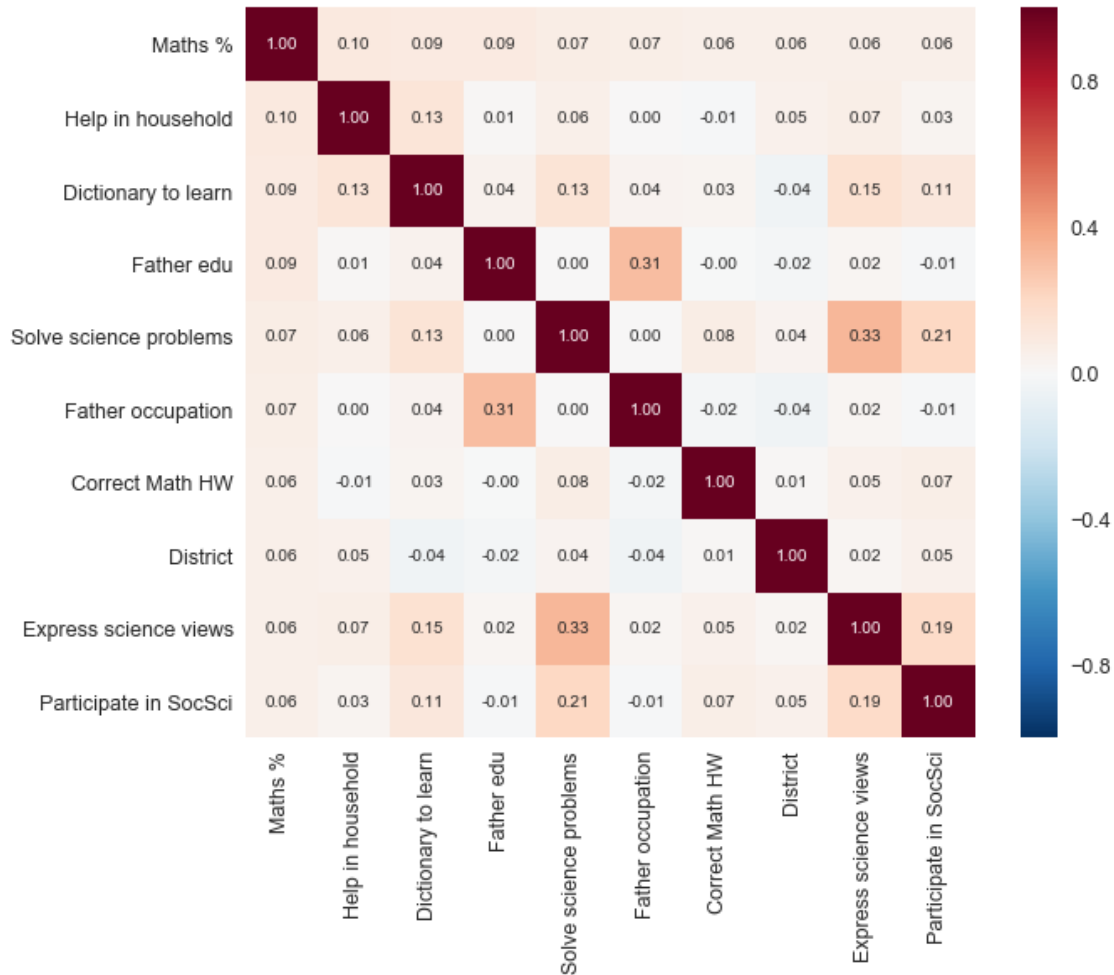
We can see now that our datasets are relatively clean. This has resulted in a rather large reduction in our dataset - each of our target variables has approximately 30% of the total samples which is less than ideal, but it's perhaps the safest way to ensure there's no tainting of our data.

1.3 Exploring Student Performance

The First Question, attempts to answer which of the features affect student performance the most. We can examine both linear and non linear correlations using various feature selection techniques. First via a correlation matrix and then using random forests for feature selection.


```
In [28]: def correlation_matrix(df, target):
    '''Helper function for correlation matrix'''
    k = 10 #number of variables for heatmap
    corrmatrix = df.corr() #Create a correlation matrix
    cols = corrmatrix.nlargest(k, target)[target].index #Find the 10 highest correlation.
    cm = np.corrcoef(df[cols].values.T)
    fig, ax = plt.subplots() #Plot this as a heatmap
    fig.set_size_inches(10,8)
    sns.set(font_scale=1.25)
    hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'s':
        yticklabels=cols.values, xticklabels=cols.values, ax=ax)
    plt.show()
```

```
In [29]: correlation_matrix(Maths, 'Maths %')
```



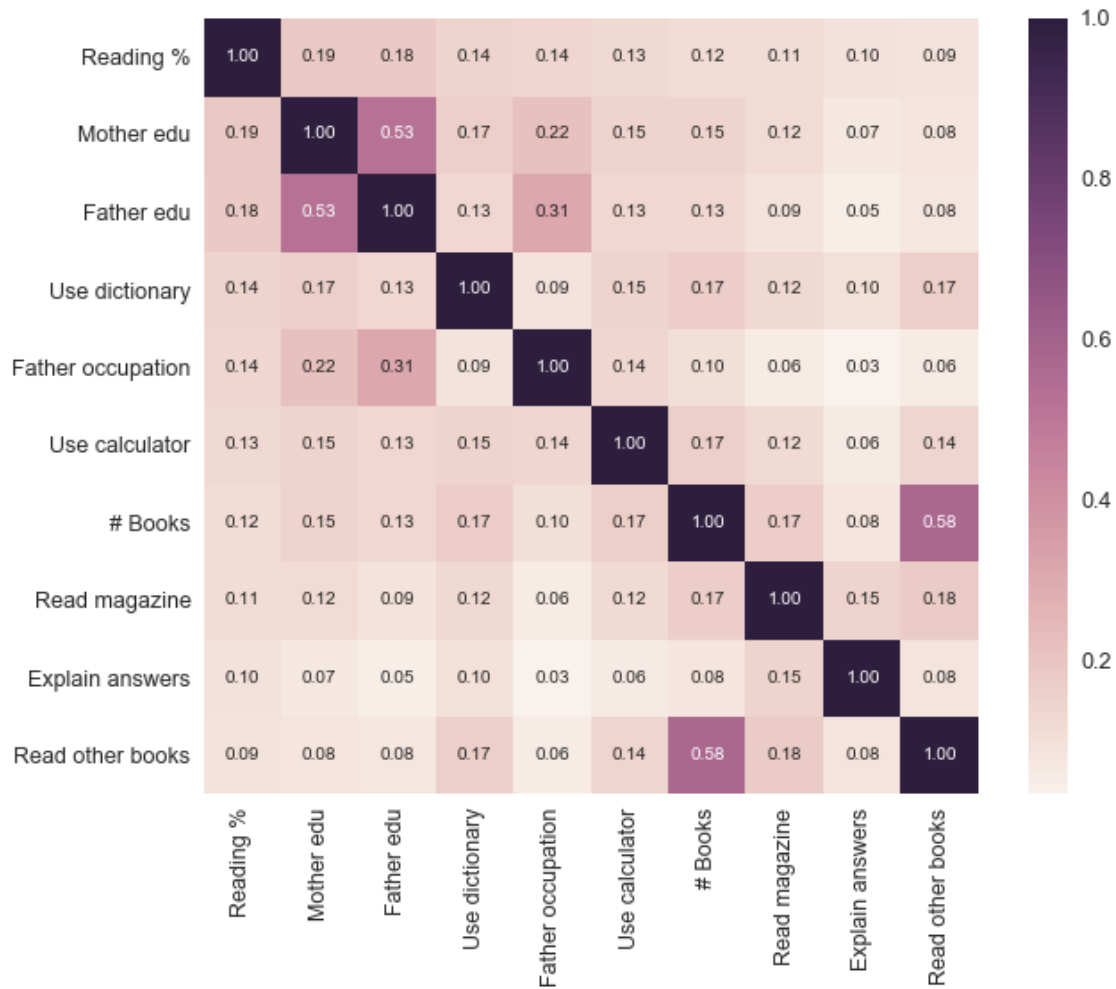
We can see from the plot that the Father's Education is quite important, as is the presence of a Dictionary, the extent to which they help in the household and the work they put into doing their homework.

```
In [30]: correlation_matrix(Science, 'Science %')
```



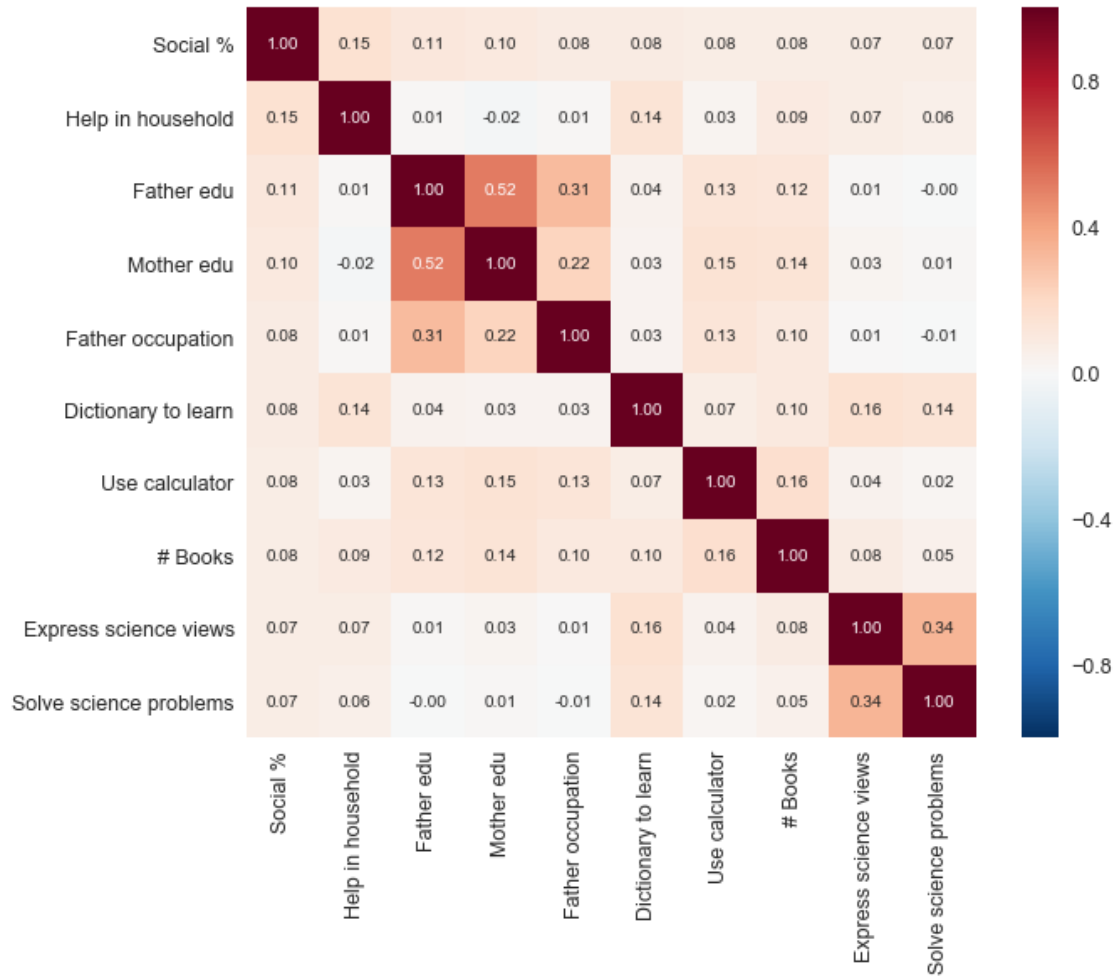
Again, Father's Education, and Mother's Education show up as important features, as does the presence of having a dictionary and the extent to which they help in the household. How much they read, also plays an important role

```
In [31]: correlation_matrix(Reading, 'Reading %')
```

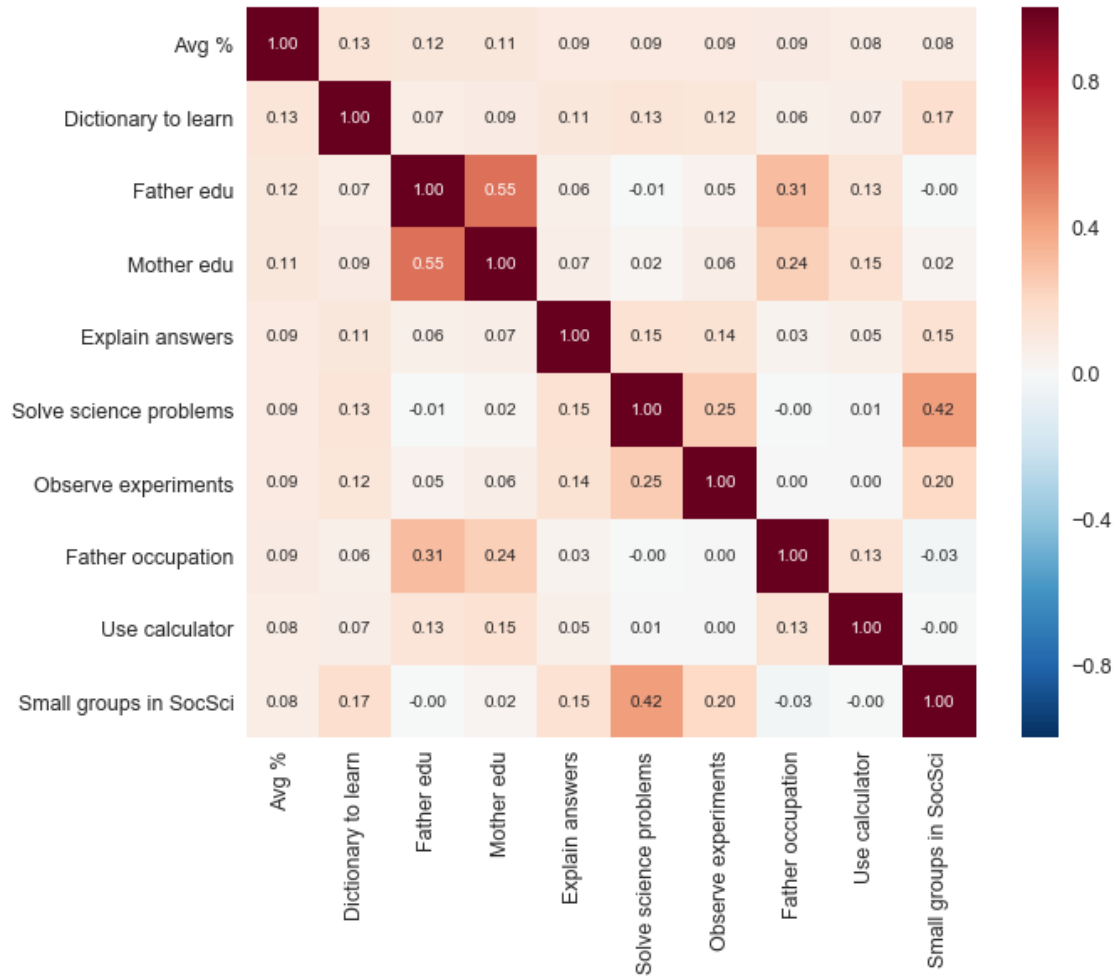



Unsurprisingly, the extent to which people read either magazines, other books etc are important towards their Reading ability. The dictionary and the parent's education is also important.

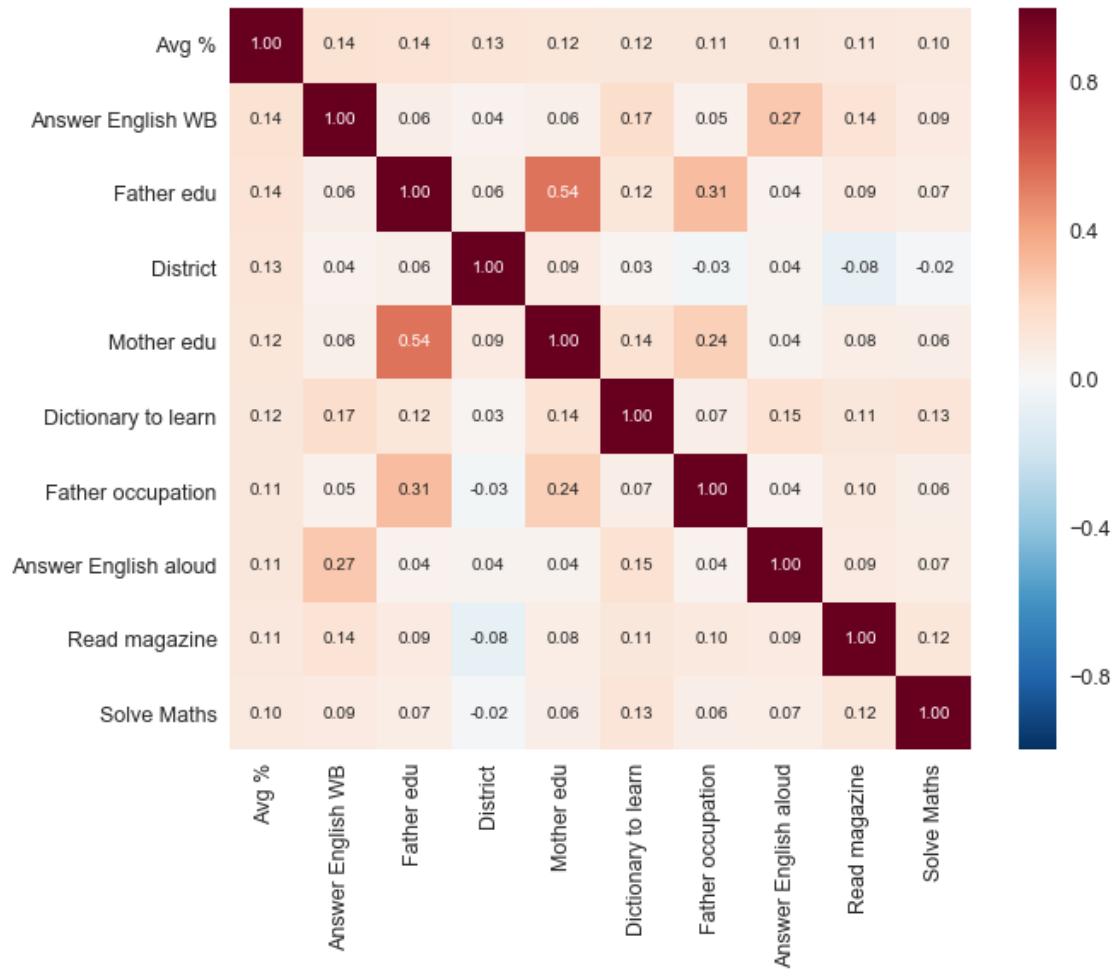
In [32]: `correlation_matrix(Social, 'Social %')`



```
In [33]: correlation_matrix(Maths_and_Science, 'Avg %')
```



```
In [34]: correlation_matrix(NullsDropped, 'Avg %')
```



Analysis In terms of a linear correlation of features, across all the datasets, the parent's education and the ability to read a dictionary is very important. How much effort the student puts in towards studies, whether it's in solving maths problems, explaining science answers, reading books and magazines, are all very important. Hard work does pay off apparently.

Now we should attempt to check for non linear correlations -

```
In [35]: from sklearn.ensemble import RandomForestRegressor
Mcopy = Maths.copy()
cat_columns = Mcopy.select_dtypes(['category']).columns
Mcopy[cat_columns] = Mcopy[cat_columns].apply(lambda x: x.cat.codes) #Encoding Categorical Variables
features = Mcopy.drop('Maths %',axis=1)
target = Mcopy['Maths %']
tree = RandomForestRegressor()
tree.fit(features,target)
sorted(zip(tree.feature_importances_,Mcopy.columns.values),reverse=True)[:10]
```

```
Out [35]: [(0.16741771951059425, 'STUID'),
(0.045822058846183962, 'State'),
```

```
(0.043747338741231746, 'District'),
(0.029811571045266709, 'Computer use'),
(0.026936486831861457, 'Father occupation'),
(0.026748832874502877, 'Father edu'),
(0.024952413896465034, 'Mother occupation'),
(0.024674615988126355, 'Library use'),
(0.024400337640052634, 'Mother edu'),
(0.022614412104687547, 'Siblings')]
```

Analysis The non linear correlations follow the linear ones, so there isn't much reason to do this for each dataset.


1.4 Girls vs Boys Across States

So, lets start by a simple metric, looking at the average of scores of boys and girls


```
In [36]: def split_girl_boy(df,target):
          boys=df.loc[df['Gender']==1].groupby('State',as_index=False)[target].mean()
          girls=df.loc[df['Gender']==2].groupby('State',as_index=False)[target].mean()
          return (boys,girls)

In [37]: (boys_maths_score,girls_maths_score) = split_girl_boy(Maths,'Maths %')
          (boys_science_score,girls_science_score) = split_girl_boy(Science,'Science %')
          (boys_reading_score,girls_reading_score) = split_girl_boy(Reading,'Reading %')
          (boys_social_score,girls_social_score) = split_girl_boy(Social,'Social %')
          (boys_mathsandscience_score,girls_mathsandscience_score) = split_girl_boy(Maths_and_Science,'Maths and Science %')
          (boys_overall_score,girls_overall_score) = split_girl_boy(NullsDropped,'Avg %')


In [38]: fig, ax = plt.subplots()
          fig.set_size_inches(15,8)
          sns.barplot(x='State',y='Maths %',hue='Gender',data=Maths, ax=ax)\
          .set_title("Maths Scores for Girls and Boys Per State")
          plt.show()
```




```
In [39]: fig, ax = plt.subplots()
fig.set_size_inches(15,8)
sns.barplot(x='State',y='Social %',hue='Gender',data=Social, ax=ax)\
.set_title("Social Science Scores per Gender per State")
plt.show()
```




```
In [40]: fig, ax = plt.subplots()
fig.set_size_inches(15,8)
sns.barplot(x='State',y='Reading %',hue='Gender',data=Reading, ax=ax)\
.set_title("Reading Scores per Gender and State")
plt.show()
```



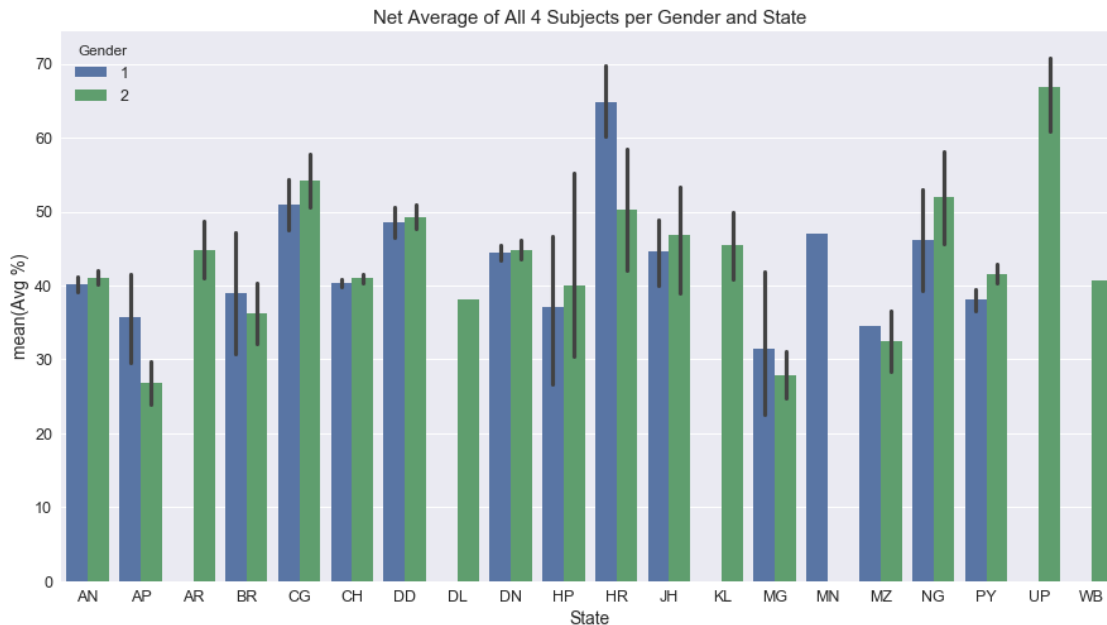
```
In [41]: fig, ax = plt.subplots()
fig.set_size_inches(15,8)
sns.barplot(x='State',y='Science %',hue='Gender',data=Science, ax=ax)\
.set_title("Science Scores per Gender and State")
plt.show()
```



```
In [42]: fig, ax = plt.subplots()
fig.set_size_inches(15,8)
sns.barplot(x='State',y='Avg %',hue='Gender',data=Maths_and_Science, ax=ax)\
.set_title("Average of Maths and Science per Gender and State")
plt.show()
```





```
In [43]: fig, ax = plt.subplots()
fig.set_size_inches(15,8)
sns.barplot(x='State',y='Avg %',hue='Gender',data=NullsDropped, ax=ax)\
.set_title("Net Average of All 4 Subjects per Gender and State")
plt.show()
```




Analysis Primarily, we can see that girls and boys have fairly comparable performance, except in certain standout cases. In maths for example, except for a few standout states, girls and boys do about the same. However, in Science, boys seem to do a lot better. Which is perhaps why Maths and Science shows the same trend as Science. Conversely girls do a lot better than boys in both reading and social science. Our data seems to back up commonly held stereotypes. As somewhat expected just taking an average of all three is such a small dataset that it doesn't seem worth to try and extrapolate information from. Contrary to most stereotypes, students from UP score very highly relative to other states. Students in Kerala perform significantly better than most other states in Reading, which lends some credence to the notion that Kerala has the highest literacy rates. Tripura and Daman and Diu have the best Science scores Overall, across all states, their average scores in Maths, Social Sciences and Science are roughly similar, and significantly lower than Reading scores. Just to get a better idea of the distribution marks over all the states, let's plot some distributions

```
In [44]: fig, ax = plt.subplots()
sns.distplot(a=boys_maths_score['Maths %'],ax=ax,label='Boys Scores')
sns.distplot(a=girls_maths_score['Maths %'],ax=ax,label='Girls Scores')
ax.legend()
plt.show()
```




```
In [45]: fig, ax = plt.subplots()
sns.distplot(a=boys_science_score['Science %'],ax=ax,label='Boys Scores')
sns.distplot(a=girls_science_score['Science %'],ax=ax,label='Girls Scores')
ax.legend()
plt.show()
```




This is probably the only dataset where girls do significantly worse in terms of performance. Just looking at the peak around the 35-40% average score for girls across states is indicative of the fact that girls don't have the best marks in science.


```
In [46]: fig, ax = plt.subplots()
sns.distplot(a=boys_reading_score['Reading %'],ax=ax,label='Boys Scores')
sns.distplot(a=girls_reading_score['Reading %'],ax=ax, label='Girls Scores')
ax.legend()
plt.show()
```



```
In [47]: fig, ax = plt.subplots()
sns.distplot(a=boys_social_score['Social %'],ax=ax,label='Boys Scores')
sns.distplot(a=girls_social_score['Social %'],ax=ax,label='Girls Scores')
ax.legend()
plt.show()
```



```
In [48]: fig, ax = plt.subplots()
sns.distplot(a=boys_mathsandscience_score['Avg %'],ax=ax,label='Boys Scores')
sns.distplot(a=girls_mathsandscience_score['Avg %'],ax=ax,label='Girls Scores')
ax.legend()
plt.show()
```



Analysis Looking at the average score of the math and science scores for girls and boys across all states, it seems like they seem to perform similarly, despite the fact that girls do relatively worse in science. Girls do significantly better in Reading, and slightly better in Social Science subjects.

1.5 Examining Performance for South Indian States

For this final question, we transform our dataset slightly. We already have some small idea about this, as in our state wise plots we could see certain states do significantly better than others in certain subjects. Let's change our levels for the state column by manufacturing another feature which aggregates whether people are in South Indian States or not. South Indian States and their respective labels in our dataset are- 1. Andaman and Nicobar (AN) 2. Andhra Pradesh (AP) 3. Karnataka (KA) 4. Kerala (KL) 5. Puducherry (PY) 6. Tamil Nadu (TN) 7. Telangana -Not in Dataset 8. Lakshadweep -Not in Dataset

Then we can use the same charts as we did previously to try and answer this question.

```
In [49]: def create_south_indian_column(df):
        south_indian_states = ['AN', 'AP', 'KA', 'KL', 'PY', 'TN']
        df['South Indian'] = [1 if x in south_indian_states else 0 for x in df['State']]
        return df
```

```
In [50]: Maths = create_south_indian_column(Maths)
        Science = create_south_indian_column(Science)
```

```


Maths_and_Science = create_south_indian_column(Maths_and_Science)
Reading = create_south_indian_column(Reading)
Social = create_south_indian_column(Social)

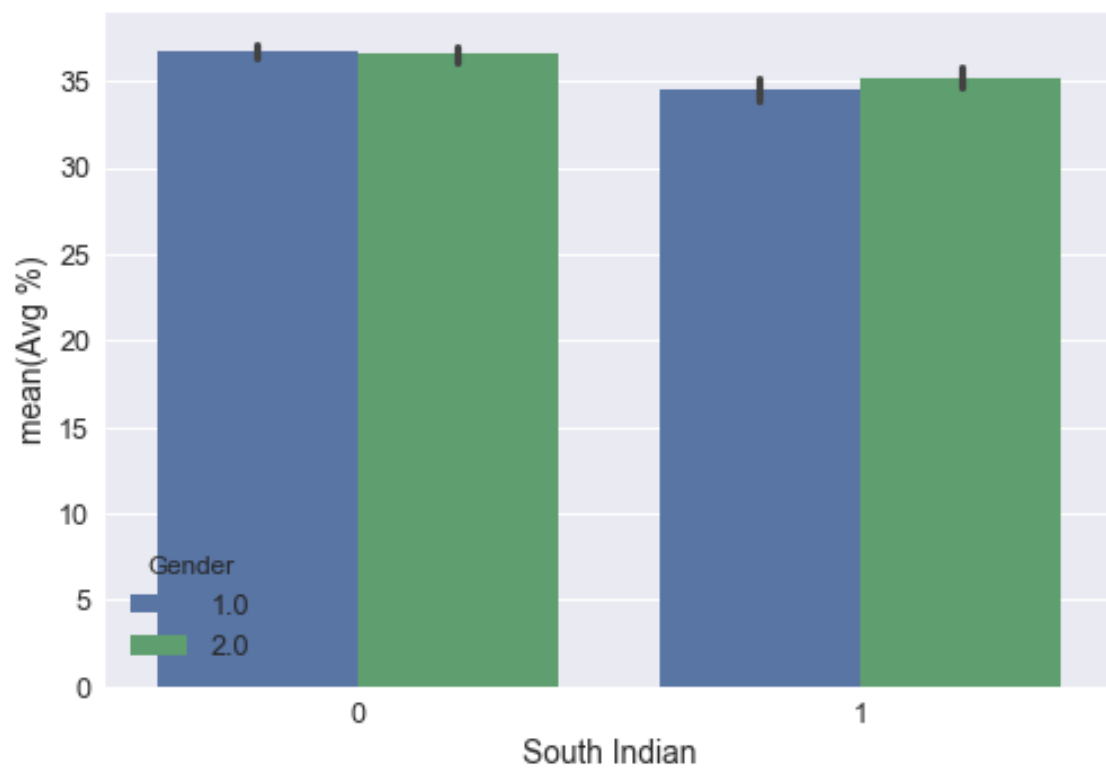
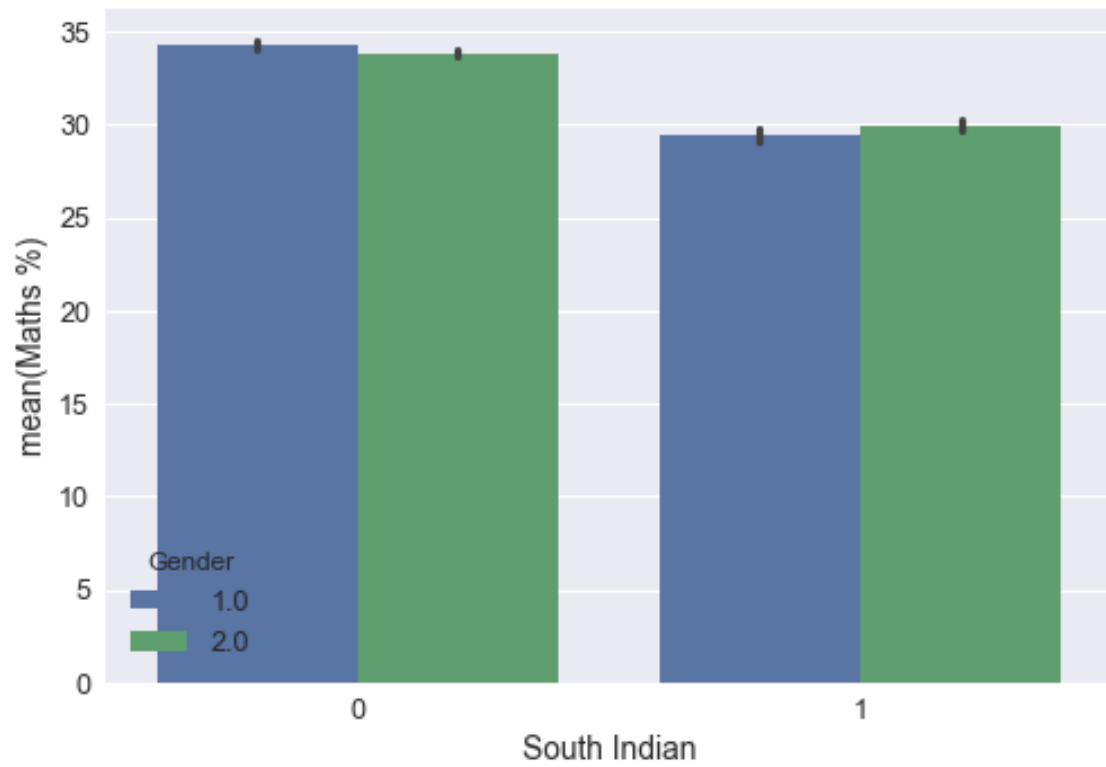
```

```


In [51]: fig, ax = plt.subplots()
sns.barplot(x='South Indian',y='Science %',hue='Gender',data=Science, ax=ax)
plt.show()
fig, ax = plt.subplots()
sns.barplot(x='South Indian',hue='Gender',y='Maths %',data=Maths, ax=ax)
plt.show()
fig, ax = plt.subplots()
sns.barplot(x='South Indian',hue='Gender',y='Avg %',data=Maths_and_Science, ax=ax)
plt.show()

```








```
In [52]: fig, ax = plt.subplots()
sns.distplot(a=Maths.loc[Maths['South Indian']==1]['Maths %'],ax=ax,label = 'South Indian')
sns.distplot(a=Maths.loc[Maths['South Indian']==0]['Maths %'],ax=ax,label='Rest of India')
ax.legend()
plt.show()
```



```
In [53]: fig, ax = plt.subplots()
sns.distplot(a=Science.loc[Science['South Indian']==1]['Science %'],ax=ax,label='South Indian')
sns.distplot(a=Science.loc[Science['South Indian']==0]['Science %'],ax=ax,label='Rest of India')
ax.legend()
plt.show()
```



```
In [54]: fig, ax = plt.subplots()
sns.distplot(a=Maths_and_Science.loc[Maths_and_Science['South Indian']==1]['Avg %'], ax=ax)
sns.distplot(a=Maths_and_Science.loc[Maths_and_Science['South Indian']==0]['Avg %'], ax=ax)
ax.legend()
plt.show()
```



From these graphs, I don't think we can conclude that students in South Indian States are extraordinarily gifted or perform significantly better in maths and science, in fact the opposite seems to be the case in maths, even if the science scores are very similar.